

# ThreadPool

# ThreadPool

- a pool of two types of threads:
  - **worker threads** - execute work item callbacks, timer expiration callbacks, and wait for registration callbacks
  - **I/O completion port thread** (IOCP threads) - execute only I/O completion callbacks
- uses background threads
- is a process-wide, single pool
- threads destroyed after some time (trimming)
- used internally by many services (WCF, Remoting, ASP.NET/ASP.NET Core), timers, TPL, PLINQ, *event-based asynchronous pattern* (EAP) and BeginXXX methods (*asynchronous programming model pattern*), even multi-tier JIT
- has a self balancing mechanism - tries to keep a "good" number of threads
- no priority support of work items - we need to write our own
- resets properties of the thread returned to the pool, including any culture that has been left behind, thread priority and the thread name (i.e., changes made with the Thread.Name property) ensure that the thread is still marked as a background thread (i.e., Thread.IsBackground is true) so that it won't hold up process exit

# ThreadPool - API

Most commonly used:

```
ThreadPool.QueueUserWorkItem(WaitCallback)
```

Sometimes used (we'll cover them too):

```
ThreadPool.GetAvailableThreads(Int32, Int32)  
ThreadPool.GetMaxThreads(Int32, Int32)  
ThreadPool.GetMinThreads(Int32, Int32)  
ThreadPool.SetMaxThreads(Int32, Int32)  
ThreadPool.SetMinThreads(Int32, Int32)
```

Much more advanced (we'll cover them only briefly):

```
ThreadPool.RegisterWaitForSingleObject(WaitHandle, WaitOrTimerCallback, Object, Int32, Boolean)  
ThreadPool.BindHandle(SafeHandle)  
ThreadPool.UnsafeQueueNativeOverlapped(NativeOverlapped*)  
ThreadPool.UnsafeQueueUserWorkItem(WaitCallback, Object)  
ThreadPool.UnsafeRegisterWaitForSingleObject(WaitHandle, WaitOrTimerCallback, Object, Int32, Boolean)  
ThreadPool.UnsafeRegisterWaitForSingleObject(WaitHandle, WaitOrTimerCallback, Object, UInt32, Boolean)
```

# ThreadPool - basic API

ThreadPool.QueueUserWorkItem

Queues a *work item* (the delegate) to be executed **at some time** by one of the ThreadPool's *worker threads*.

- does not guarantee to be executed immediately
- there is no direct way to wait for it or get a result ("*fire & forget*")
- has some caveats (we'll return to that...)

```
public static void Main()
{
    ThreadPool.QueueUserWorkItem(DoWork);
    Console.ReadLine();
}

static void DoWork(object? state)
{
    // Do something...
}
```

# ThreadPool - exception handling

- if work item throws an exception, it is unhandled - kaput!
- we can use `AppDomain.CurrentDomain.UnhandledException` to log & save the work (but not to prevent the app closing)

```
public static void Main()
{
    AppDomain.CurrentDomain.UnhandledException += CurrentDomainOnUnhandledException;
    ThreadPool.QueueUserWorkItem(DoWork);
    Console.ReadLine();
}

private static void CurrentDomainOnUnhandledException(object sender, UnhandledExceptionEventArgs e)
{
    Console.WriteLine($"{(e.IsTerminating ? "Terminating" : "Non-terminating")} fault: {e.ExceptionObject}");
}

static void DoWork(object? state)
{
    throw new NullReferenceException();
}
```

# ThreadPool - termination

- again: threads in the thread pool are **background threads**
- .NET process doesn't wait for work items (including those queued)
- again: cooperate by observing passed `CancellationToken.IsCancellationRequested`

# ThreadPool - termination

Want to make sure processing will end?

- consider using different mechanism - separate thread(s) with some synchronization
- you *theoretically* can change thread pool's thread to be a foreground thread

# ThreadPool - termination

## Guarding by making threads foreground

```
public static void Main()
{
    ThreadPool.QueueUserWorkItem(DoWork);
    // main thread ends but will DoWork __if already queued__
}

static void DoWork(object? state)
{
    Thread.CurrentThread.IsBackground = false;
    // ...
}
```

It works because `IsBackground` is reset when the thread ends executing work item (so no other work items will be influenced).

Obvious two problems:

- the main thread (and the whole app) may end before queued work item was started
  - we could overcome that by observing `ThreadPool.PendingWorkItemCount` from another foreground thread
- even if enqueued, the `DoWork` may be terminated by app close before setting `IsBackground` flag
  - it needs some additional scheduling/synchronization mechanism



# ThreadPool - managing the pool

- the thread pool starts with some default number of threads in it (zero?).
- then it needs to add/remove them as needed:
  - tries to keep the number of threads low, tunes to the workload, and **avoids creating new threads as long as possible**
  - too few threads - if many work items waiting, we waste CPU
  - too many threads - context switching/cache misses and CPU trashing (a lot of waits ended at the same time)
- in fact - what is the "wait ratio" of work items?
- thread pool maintains a minimum and a maximum number of threads (separately for worker and IOCP)

# ThreadPool - managing the pool

## maximum number of threads

- means a real maximum
- there will be no more worker/IOCP threads created than this value
- if work items count exceeded - the work item will be **queued** (until one worker threads becomes available)

# ThreadPool - managing the pool

## maximum number of threads

- default value for worker threads:
  - .NET Framework 2.0 - 25 per core
  - .NET Framework 3.5 - 250 per core
  - .NET Framework 4.0 in a 32-bit environment - 1,023
  - **.NET Core x64**, .NET Framework 4.0+ - 32,768
  - .NET Core x86 -  $X$ 
    - $X = (\text{half of 32-bit address space}) / (\text{default stack size})$
    - default stack size taken from EXE (typically 1.5MB on Windows) or 256 kB (assumed on Linux)
    - thus typically  $X$  is 682 (Windows)
- default value for IOCP threads:
  - 1000

[More details in win32threadpool.cpp for brave ones](#)

# ThreadPool - managing the pool

## minimum number of threads

- does not mean a real minimum!
- there may be fewer threads than this value (including 0)
- if there are fewer current threads than a minimum - create threads **immediately** as needed
- if there are more current threads than a minimum - create every new thread **at maximum 0.5 second intervals**

# ThreadPool - managing the pool

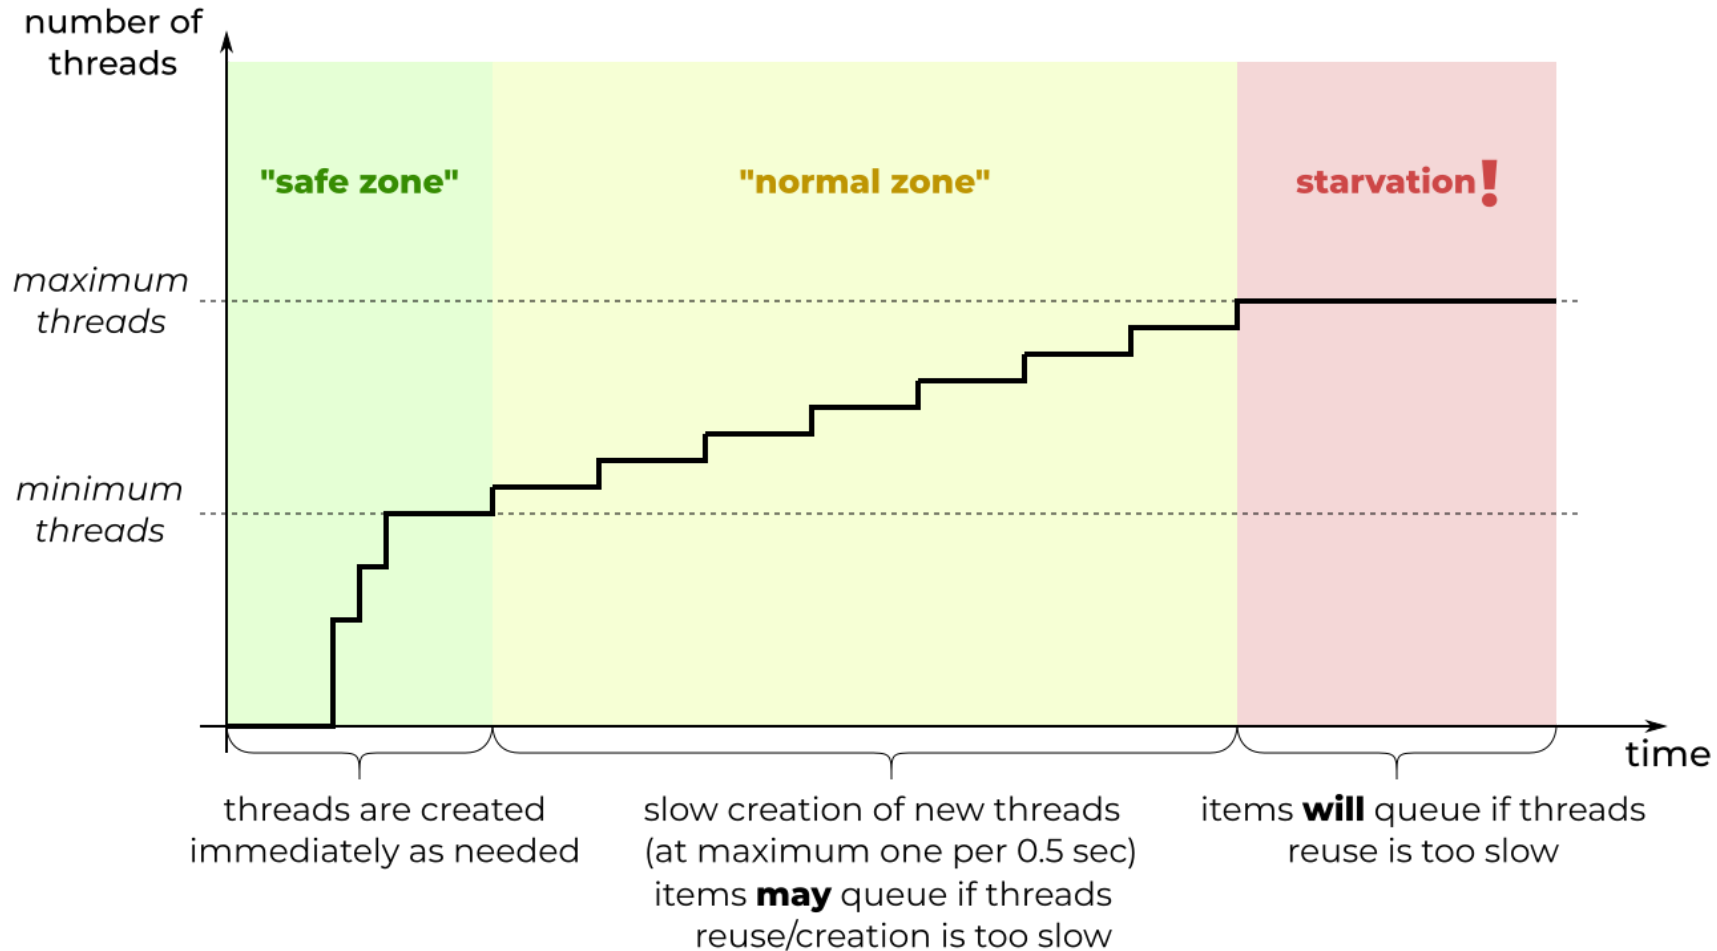
## minimum number of threads

- default value for worker & IOCP threads: number of processors (logical cores)
- raising it may improve concurrency and better "burst response" in case of many blocked threads

IOCP threads throttled by IOCP/Windows

# ThreadPool - managing the pool

*Note: Assuming a lot of work items is arriving so threads reuse is not enough*



# ThreadPool - managing the pool

Two main algorithms for managing threads in "normal zone":

- a starvation-avoidance mechanism - adding worker threads if it sees no progress being made on queued items
- a *hill-climbing heuristic* - maximize throughput while using as few threads as possible (injecting/removing threads)

An opportunity to inject threads (whichever is shorter):

- every time a work item completes, or
- **at maximum 0.5 second intervals** (if *minimum* exceeded)

# ThreadPool - managing the pool

Hill-climbing described in ["Optimizing Concurrency Levels in the .NET ThreadPool"](#) paper published by Microsoft.

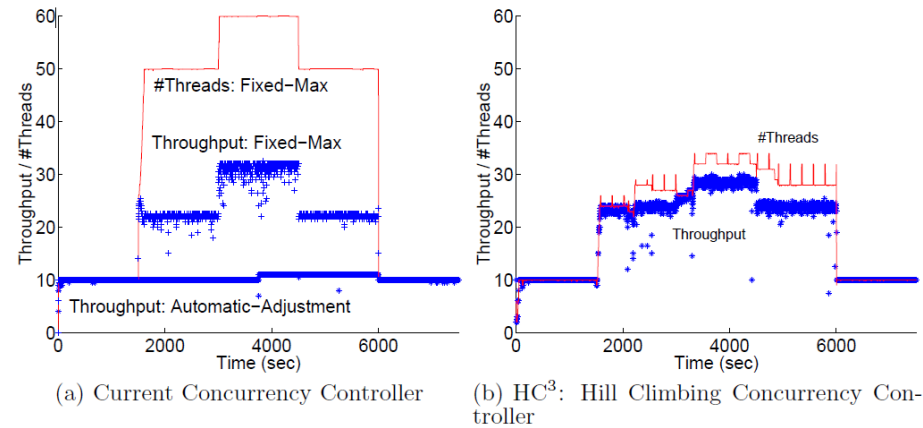


Figure 3: Performance of Concurrency Controllers for the CLR ThreadPool for a Dynamic Workload

In general: *"adapt quickly, but not too quickly"*, basing on:

- current thread count,
- current number of items completed/time,
- queueing rate,
- CPU utilization



# ThreadPool - managing the pool

## Summary

- in .NET Core it is pretty hard to have a real **thread pool starvation**
  - exceeding 32,768 work items...
- but we may lead to "soft" thread pool starvation
  - if we have more than *minimum* and there is a "rush" of requests, new threads will be created very slowly
  - work items will queue due to not enough worker threads
  - hill-climbing tries to accommodate that (and eventually will)
  - such state is mostly dangerous for unexpected "peaks" of work items
- it may influence... everything: HTTP request processing, async calls, timers, even JIT

# ThreadPool - configuration

## .NET Framework

Configuration element in machine.config (yes, system-wide):

```
<processModel autoConfig="false" minIoThreads="250" />
```

Remember: *"the value specified in this configuration element is a per-core setting"*

## .NET Core

[appname].runtimeconfig.json

```
{
  "runtimeOptions": {
    "configProperties": {
      "System.GC.Server": true,
      "System.GC.Concurrent": true,
      "System.Threading.ThreadPool.MinThreads": 4,
      "System.Threading.ThreadPool.MaxThreads": 8
    },
    ...
  }
}
```

# ThreadPool - configuration

## Programatically

```
ThreadPool.SetMinThreads(int workerThreads, int completionPortThreads)  
ThreadPool.SetMaxThreads(int workerThreads, int completionPortThreads)
```

Sidenote: You cannot set the maximum number of worker threads or I/O completion threads to a number smaller than the number of processors on the computer.

## **Demo:** ThreadPool **use with different API**

# ThreadPool - waits

**Registered wait** - a special, single *wait thread* created for every 63 objects registered - This thread manages waiting on objects and queuing the callbacks to run on one of the thread pool's worker threads when an object is signaled

# ThreadPool - waits

ThreadPool.RegisterWaitForSingleObject - asynchronously waiting on a wait handle, instead of a synchronous .WaitOne() that blocks a thread

```
class CallbackInfo
{
    public RegisteredWaitHandle Handle;
}

class RegisteredWaitDemo
{
    public static void Main(string[] args)
    {
        AutoResetEvent evt = new AutoResetEvent(false);
        var callbackInfo = new CallbackInfo();
        callbackInfo.Handle = ThreadPool.RegisterWaitForSingleObject(evt, CallBack, state: callbackInfo,
            timeout: TimeSpan.FromSeconds(3.0), executeOnlyOnce: true);

        Thread.Sleep(2000); // Start long running operation
        evt.Set(); // ... and signal it has finished
        // Dispose AutoResetEvent or reuse for other notifications
    }

    private static void CallBack(object? state, bool timedOut)
    {
        var callbackInfo = state! as CallbackInfo;
        callbackInfo?.Handle?.Unregister(null); // Good practice to unregister
    }
}
```

# ThreadPool - waits

Executes action on a thread pool thread after sleep milliseconds - consumes only *wait thread* for waiting:

```
public static void ExecuteAfterWait<T>(int sleep, Action<T> action, T arg)
{
    AutoResetEvent evt = new AutoResetEvent(false);
    RegisteredWaitHandle? handle = null;
    handle = ThreadPool.RegisterWaitForSingleObject(evt, (state, timedOut) =>
    {
        handle!.Unregister(evt);
        action(arg);
    }, null, sleep, true);
}
```

# ThreadPool internals - task queues (since .NET 4)

A *work item* will land in one of two queues:

- global task queue
  - if the item has been enqueued from a non-thread pool thread
  - if we use `ThreadPool.QueueUserWorkItem/ThreadPool.UnsafeQueueUserWorkItem`
  - if we use `TaskCreationOptions.PreferFairness` flag with `Task.Factory.StartNew`
  - if we call `Task.Yield` on the default task scheduler
- local task queues for each worker thread
  - if the item has been enqueued from a thread pool thread



# ThreadPool internals - dequeuing

Worker thread looks:

- into its local queue (LIFO order)
  - LIFO to preserve cache/data locality
- if empty, into global queue (FIFO order)
- if empty, into local queues of other threads ("*work stealing*", FIFO order)

# ThreadPool internals - dequeuing

[\*".NET Threadpool starvation, and how queuing makes it worse"\*](#) by Criteo - great article about example deadlock scenario because of that:

# ThreadPool internals - dequeuing

[\*".NET Threadpool starvation, and how queuing makes it worse"\*](#) by Criteo - great article about example deadlock scenario because of that:

1. a constant burst of requests comes

# ThreadPool internals - dequeuing

[".NET Threadpool starvation, and how queuing makes it worse"](#) by Criteo - great article about example deadlock scenario because of that:

1. a constant burst of requests comes
2. thread pool becomes starved (*softly*)

# ThreadPool internals - dequeuing

[".NET Threadpool starvation, and how queuing makes it worse"](#) by Criteo - great article about example deadlock scenario because of that:

1. a constant burst of requests comes
2. thread pool becomes starved (*softly*)
3. enqueue 5 items per second **to the global queue** (because of burst)

# ThreadPool internals - dequeuing

[".NET Threadpool starvation, and how queuing makes it worse"](#) by Criteo - great article about example deadlock scenario because of that:

1. a constant burst of requests comes
2. thread pool becomes starved (*softly*)
3. enqueue 5 items per second **to the global queue** (because of burst)
4. each of such items enqueues one more item into the local queue **and waits for it**

# ThreadPool internals - dequeuing

[".NET Threadpool starvation, and how queuing makes it worse"](#) by Criteo - great article about example deadlock scenario because of that:

1. a constant burst of requests comes
2. thread pool becomes starved (*softly*)
3. enqueue 5 items per second **to the global queue** (because of burst)
4. each of such items enqueues one more item into the local queue **and waits for it**
5. when a new thread is created (because of starvation), it looks into:
  - its local queue - but it is empty, as it was just created
  - the global queue - which is constantly growing

# ThreadPool internals - dequeuing

[".NET Threadpool starvation, and how queuing makes it worse"](#) by Criteo - great article about example deadlock scenario because of that:

1. a constant burst of requests comes
2. thread pool becomes starved (*softly*)
3. enqueue 5 items per second **to the global queue** (because of burst)
4. each of such items enqueues one more item into the local queue **and waits for it**
5. when a new thread is created (because of starvation), it looks into:
  - its local queue - but it is empty, as it was just created
  - the global queue - which is constantly growing
6. enqueueing to the global queue (**5/sec**) is faster than the threadpool grow (**0.5/sec**) so we never recover - locally queued tasks are not processed as newly created threads are taking tasks from the global queue. So worker threads are blocked waiting...



# ThreadPool internals - dequeuing

["NET Threadpool starvation, and how queuing makes it worse"](#) by Criteo - great article about example deadlock scenario because of that:

1. a constant burst of requests comes
2. thread pool becomes starved (*softly*)
3. enqueue 5 items per second **to the global queue** (because of burst)
4. each of such items enqueues one more item into the local queue **and waits for it**
5. when a new thread is created (because of starvation), it looks into:
  - its local queue - but it is empty, as it was just created
  - the global queue - which is constantly growing
6. enqueueing to the global queue (**5/sec**) is faster than the threadpool grow (**0.5/sec**) so we never recover - locally queued tasks are not processed as newly created threads are taking tasks from the global queue. So worker threads are blocked waiting...


## Solution:

- change step 3 to enqueue **to the local queue**
- when a new thread is created (because of starvation), it looks into:
  - its local queue - but it is empty, as it was just created
  - the global queue - but it is empty, as we not enqueue there anything
  - steals the work from the other local queue
- never, ever wait synchronously!

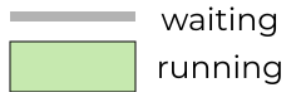
# ThreadPool internals - I/O operations

- we issue an **overlapped I/O**
  - it may complete immediately
  - or, returns a status code indicating the I/O operation is still pending.
- I/O completion port - to observe the result
  - it may represent many simultaneous ongoing I/O operations
  - one or many threads may wait for the IOCP
  - then it executes some callback with the incoming data
- in case of .NET ThreadPool:
  - one of IOCP threads will observe the change
  - mark the task complete
  - the continuation is queued to the target SynchronizationContext or the thread pool
- IOCP threads are blocked, yes - they block on the IOCP itself.
  - one IOCP thread can handle a huge number of IOCP notifications,

# ThreadPool internals - I/O operations



- 1** IOCP threads are blocked for notification on assigned IOCP
- 2** worker thread "opens" device and "binds" it to IOCP
- 3** worker thread starts overlapped I/O (providing **callback** to execute)
- 4** I/O operation completes
- 5** IOCP is signalled
- 6** one of the blocked IOCP threads is signalled and executes **callback** (it may delegate further work to a worker thread)
- 7** one of the worker threads is signalled



# ThreadPool - monitoring

Programatically:

```
int workerThreads = 0;  
int completionPortThreads = 0;  
ThreadPool.GetAvailableThreads(out workerThreads, out completionPortThreads);  
  
ThreadPool.GetMinThreads(out int workerThreads, out int completionPortThreads);  
ThreadPool.GetMaxThreads(out int workerThreads, out int completionPortThreads);
```

# ThreadPool - monitoring

dotnet counters:

[System.Runtime]	
% Time in GC since last GC (%)	0
Allocation Rate / 1 sec (B)	0
CPU Usage (%)	0
Exception Count / 1 sec	0
GC Heap Size (MB)	1
Gen 0 GC Count / 60 sec	0
Gen 0 Size (B)	0
Gen 1 GC Count / 60 sec	0
Gen 1 Size (B)	0
Gen 2 GC Count / 60 sec	0
Gen 2 Size (B)	0
LOH Size (B)	0
Monitor Lock Contention Count / 1 sec	0
Number of Active Timers	0
Number of Assemblies Loaded	10
ThreadPool Completed Work Item Count / 1 sec	0
ThreadPool Queue Length	1,978
ThreadPool Thread Count	25
Working Set (MB)	21

or... ASP.NET 4.0.x/Request Queued performance counter in case of ASP.NET Framework

## ThreadPool - monitoring

## PerfView (standard *.NET* session) or dotnet-trace (opened in PerfView)

[illegible]

# ThreadPool - monitoring

Dump analysis with the help of SOS threadpool command:

```
> dotnet dump collect -p 29924
Writing full to dump_20200515_131534.dmp
Complete

> dotnet dump analyze .\dump_20200515_131534.dmp
Loading core dump: .\dump_20200515_131534.dmp ...
Ready to process analysis commands. Type 'help' to list available commands or 'help [command]' to get detailed
help on a command.
Type 'quit' or 'exit' to exit the session.
> threadpool
CPU utilization: 17%
Worker Thread: Total: 8 Running: 8 Idle: 0 MaxLimit: 8 MinLimit: 8
Work Request in Queue: 1
    AsyncTimerCallbackCompletion TimerInfo@000001E91E19FE50
-----
Number of Timers: 0
-----
Completion Port Thread: Total: 0 Free: 0 MaxFree: 16 CurrentLimit: 0 MaxLimit: 1000 MinLimit: 8
```

# Scenario - StackExchange.Redis

[Are you seeing a high number of busyio or busyworker threads in the timeout exception?](#)

StackExchange.Redis:

- has its own dedicated thread-pool (to have better control) - serves most calls
- fall-backs to .NET ThreadPool if dedicated is being full

*"Note that **StackExchange.Redis can hit timeouts** if either the IOCP threads or the worker threads (.NET global thread-pool, or the dedicated thread-pool) become saturated without the ability to grow."*

*"Basically, if you're hitting the global thread pool (rather than the dedicated StackExchange.Redis thread-pool) it means that when the number of Busy threads is greater than Min threads, **you are likely paying a 500ms delay** before network traffic is processed by the application."*

*"If we look at an example error message from StackExchange.Redis 2.0, you will see that it now prints ThreadPool statistics (see IOCP and WORKER details below)"*

```
Timeout performing GET MyKey (1000ms), inst: 2, qs: 6, in: 0, mgr: 9 of 10 available,  
IOCP: (Busy=6,Free=994,Min=4,Max=1000),  
WORKER: (Busy=3,Free=997,Min=4,Max=1000)
```

*"A good starting place is 200 or 300, then test and tweak as needed."*



# ThreadPool - advanced configuration

[Advanced CLR configuration knobs](#) about thread pool:

- `ThreadPool_DebugBreakOnWorkerStarvation` - breaks into the debugger if the ThreadPool detects work queue starvation (default 0)
- `ThreadPool_DisableStarvationDetection` - disables the ThreadPool feature that forces new threads to be added when workitems run for too long (default 0)
- `ThreadPool_EnableWorkerTracking` - enables **extra expensive** tracking of how many workers threads are working simultaneously (default 0)
- `ThreadPool_ForceMaxWorkerThreads` - overrides the `MaxThreads` setting for the ThreadPool worker pool (default 0)
- `ThreadPool_ForceMinWorkerThreads` - overrides the `MinThreads` setting for the ThreadPool worker pool (default 0)
- `ThreadPool_UnfairSemaphoreSpinLimit` - maximum number of spins per processor a thread pool worker thread performs before waiting for work (default 0x32)
- `ThreadPoolTickCountAdjustment` - internal setting used only in the Debug build of the runtime
- `HillClimbing_` - *a bunch of settings* for Hill climbing\* algorithm

# ThreadPool - advanced configuration

ThreadPool\_EnableWorkerTracking **example** - *"enables extra expensive tracking of how many workers threads are working simultaneously"*

Event Name	Time MSec	Process Name	Count
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	30,671.967	StandaloneThreadPoolConsoleApp (24324)	7
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	31,170.799	StandaloneThreadPoolConsoleApp (24324)	7
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	31,672.373	StandaloneThreadPoolConsoleApp (24324)	8
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	32,171.680	StandaloneThreadPoolConsoleApp (24324)	7
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	32,671.572	StandaloneThreadPoolConsoleApp (24324)	7
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	33,171.331	StandaloneThreadPoolConsoleApp (24324)	7
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	33,671.010	StandaloneThreadPoolConsoleApp (24324)	6
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	34,171.363	StandaloneThreadPoolConsoleApp (24324)	6
Microsoft-Windows-DotNETRuntime/ThreadPoolWorkingThreadCount/Start	34,672.399	StandaloneThreadPoolConsoleApp (24324)	6

# ThreadPool internals - I/O operations (cont.)

I/O completion port is exposed through the .NET [ThreadPool.BindHandle\(SafeHandle handle\)](#) method:

- A `SafeHandle` that holds the operating system handle. The handle must have been opened for *overlapped I/O*
- you can use it to serve high-throughput I/O devices, use overlapped I/O with completion ports
- you can bind a given handle to be served by IOCP Threads

## **Demo:** ThreadPool.BindHandle **usage**

# ThreadPool - "unsafe" API

`ThreadPool.UnsafeQueueUserWorkItem()`

ExecutionContext which includes various security information (impersonation information established for the thread), is **NOT captured** at the time of the call (on the queuing thread) and so, then **NOT used** when invoking the callback on the thread pool. But **faster**.

`ThreadPool.UnsafeQueueNativeOverlapped()`

Queues an overlapped I/O operation for execution (on IOCP threads pool).

`ThreadPool.UnsafeRegisterWaitForSingleObject()`

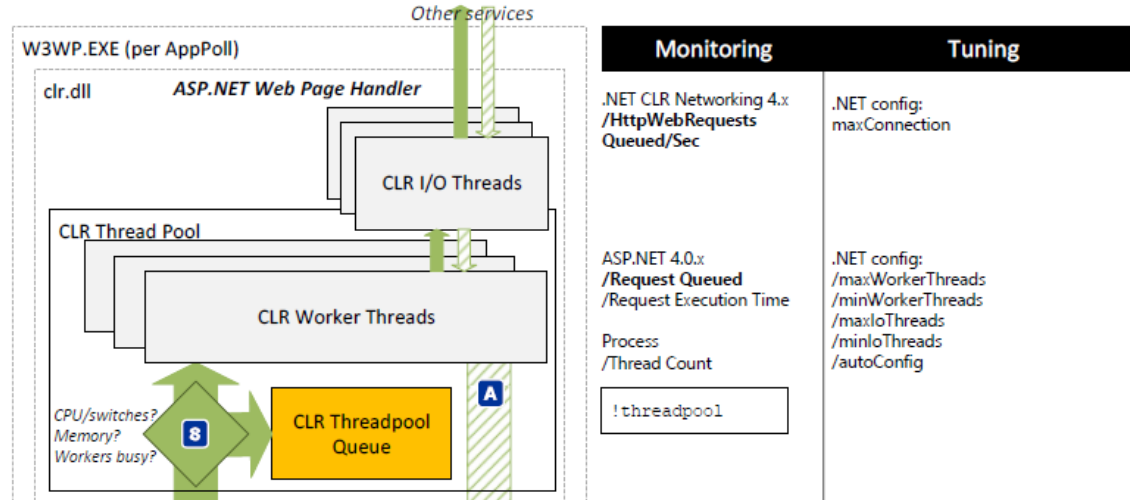
Unsafe version of RegisterWaitForSingleObject, so no ExecutionContext capturing.

# ThreadPool - Materials

- <https://dschenkelman.github.io/2013/10/29/asynchronous-io-in-c-io-completion-ports/>
- [Understanding the Windows I/O System](#)
- [Multithreaded Asynchronous I/O & I/O Completion Ports](#)
- [The CLR Thread Pool 'Thread Injection' Algorithm](#)
- [Concurrency - Throttling Concurrency in the CLR 4.0 ThreadPool](#)
- [New and Improved CLR 4 Thread Pool Engine](#)
- [Diagnosing .NET Core ThreadPool Starvation with PerfView](#)

# IIS hosting

## Kestrel hosting





# Kestrel hosting

# Kestrel hosting

- asynchronous I/O initially based on libuv library (node.js) - single-threaded even looping model
- all other work done by worker threads
- *"Kestrel also supports running multiple event loops, thereby making it a more robust choice than Node.js. The number of event loops that are used depends on the number of logical processors on the machine"*

# Kestrel hosting - ASP.NET Core 1.x

```
WebHost.CreateDefaultBuilder(args)
    .UseLibuv(opts => opts.ThreadCount = 4)
    .UseStartup<Startup>();
```

# Kestrel hosting - ASP.NET Core 2.1+

- libuv replaced by **managed sockets**
  - the Linux transport makes use of non-blocking sockets and *epoll*
  - *"The big debate we usually have around running code on epoll threads is blocking user code. As a result, we dispatch from a small number of epoll threads to the thread pool (which grows and supports work stealing)"* - David Fowler
- fall-back possible with `Microsoft.AspNetCore.Server.Kestrel.Transport.Libuv` NuGet package

# Kestrel hosting - ASP.NET Core 2.1+

```
Host.CreateDefaultBuilder(args)
    .ConfigureKestrel((context, options) =>
    {
        options.Limits.MaxConcurrentConnections = 100;
        options.Limits.MaxConcurrentUpgradedConnections = 100; // limit for connections that have been upgraded
                                                                // from HTTP or HTTPS to another protocol (for
                                                                // example, on a WebSockets request)
    })
```

appsettings.json Or appsettings.{Environment}.json:

```
{
  "Kestrel": {
    "Limits": {
      "MaxConcurrentConnections": 100,
      "MaxConcurrentUpgradedConnections": 100
    },
    ...
  }
}
```

# Kestrel hosting - ASP.NET Core 2.1+

`SocketTransportOptions.IOQueueCount` - the number of I/O queues used to process requests. Set to 0 to directly schedule I/O to the `ThreadPool`. Defaults to `ProcessorCount` rounded down and clamped between 1 and 16.

```
public static IHostBuilder CreateHostBuilder(string[] args) =>
    Host.CreateDefaultBuilder(args)
        .ConfigureWebHostDefaults(webBuilder =>
        {
            webBuilder
                .UseSockets(options => options.IOQueueCount = 0)
                .UseStartup<Startup>();
        });
```

# Kestrel hosting - custom

<https://github.com/redhat-developer/kestrel-linux-transport> - custom Transport abstraction implementation for Kestrel, uses `SO_REUSEPORT` and *POSIX AIO* (asynchronous I/O) interface

- reference `RedHat.AspNetCore.Server.Kestrel.Transport.Linux` package
- configure `WebHost`:

```
public static IWebHost BuildWebHost(string[] args) =>
    Host.CreateDefaultBuilder(args)
        .ConfigureWebHostDefaults(webBuilder =>
        {
            webBuilder
                .UseLinuxTransport()
                .UseStartup<Startup>();
        });
```

# Kestrel hosting - future

[.NET 5 on Linux: AIO, io\\_uring, and thread pool changes](#)





# **ThreadPool - asynchronous delegates**

# ThreadPool - asynchronous delegates

- `ThreadPool.QueueUserWorkItem` doesn't expose any direct way to return a result
- *asynchronous delegate* allows any number of typed arguments to be passed in both directions
- `BeginInvoke` to start and then `EndInvoke`:
  - waits for the asynchronous delegate to execute and finish on the thread pool,
  - it receives the return value
  - it (re)throws an unhandled exception

```
Func<T1, T2> method = DoWork;  
IAsyncResult handle = method.BeginInvoke(args);  
// async part  
T2 result = method.EndInvoke(handle);  
...
```

or

```
Func<T1, T2> method = DoWork;  
method.BeginInvoke(args, OnDone, method);  
// async part  
...  
static void OnDone<T1, T2>(IAsyncResult handle)  
{  
    var target = (Func<T1, T2>)handle.AsyncState;  
    T2 result = target.EndInvoke(handle);  
    ...  
}
```

**ThreadPool - asynchronous delegates**

**No-no, legacy!**

# **Asynchronous Programming Model (APM)**

# Asynchronous Programming Model (APM)

- pair of Begin.../End... methods, for example:

```
public IAsyncResult BeginRead(byte[] buffer, int offset,  
                             int count, AsyncCallback callback,  
                             object state);  
public int EndRead(IAsyncResult asyncResult); // waits for result
```

- IAsyncResult:
  - IsCompleted
  - CompletedSynchronously
  - AsyncWaitHandle
  - AsyncState
- accessing results:
  - pooling IsCompleted and then calling End...
  - waiting on AsyncWaitHandle (with timeout), i.e. with ThreadPool.RegisterWaitForSingleObject
  - using delegate void AsyncCallback(IAsyncResult iar)
  - Task<int>.Factory.FromAsync(fs.BeginRead, fs.EndRead, buffer, 0, buffer.Length, null)

# **Event Asynchronous Pattern (EAP)**

# Event Asynchronous Pattern (EAP)

- asynchronous operation start and C# event handler for results
- implicit multi-threading
- implicit synchronization with proper context (UI)
- very rare: WebClient, BackgroundWorker and exotic: SoundPlayer, System.Windows.Forms.PictureBox.LoadAsync
- those ...Async methods are not our beloved async, i.e. WebClient.DownloadStringAsync and WebClient.DownloadStringAsyncTask

# Event Asynchronous Pattern (EAP)

WebClient

```
public static void Main(string[] args)
{
    WebClient client = new WebClient();
    client.DownloadStringCompleted += OnDownloadCompleted;
    client.DownloadStringAsync(new Uri(@"https://google.com"));
    //client.CancelAsync();
}

private static void OnDownloadCompleted(object sender,
                                       DownloadStringCompletedEventArgs args)
{
    // UI is safe here
    if (args.Cancelled)
        Console.WriteLine("Canceled");
    else if (args.Error != null)
        Console.WriteLine("Exception: " + args.Error.Message);
    else
        Console.WriteLine(args.Result);
}
```



# Event Asynchronous Pattern (EAP)

BackgroundWorker

General-purpose "worker" with cancellation, progress, UI thread sync - based on the thread pool.

```
public static void Run(string[] args)
{
    var worker = new BackgroundWorker();
    worker.DoWork += WorkerOnDoWork;
    worker.ProgressChanged += WorkerOnProgressChanged;
    worker.RunWorkerCompleted += WorkerOnRunWorkerCompleted;
    worker.RunWorkerAsync(worker);
    // Dispose or reuse the worker
}

private static void WorkerOnRunWorkerCompleted(object sender, RunWorkerCompletedEventArgs e)
{
}

private static void WorkerOnProgressChanged(object sender, ProgressChangedEventArgs e)
{
}

private static void WorkerOnDoWork(object sender, DoWorkEventArgs e)
{
    BackgroundWorker worker = e.Argument as BackgroundWorker;
    // Use worker.CancellationPending, worker.ReportProgress(int) and e.Result
}
```

# Exercise

## .\Module02-AsyncBasics\Exercises

Implement methods from `ThreadPoolExercises.Core` project to pass tests from `ThreadPoolExercises.Tests`. You can use `ThreadPoolExercises` as a playground.

In the end, use `ThreadPoolExercises.Benchmarks` to benchmark implemented methods.